

Week 7

# LayerNorm Optimization

- Implementation 1:
  - Loop 1: mean
  - Loop 2: variance
  - Loop 3: scaling
- Implementation 2:
  - Loop 1: mean & variance
    - Use current mean to calculate variance
  - Loop 2: scaling

```
def forward(self, x):
    row = x.size()[-2]
    col = x.size()[-1]

    # Initialize mean and variance with zero
    # mean = x[0, :, 0] * 0
    # var = x[0, :, 0] * 0

    mean = x[:, 0] * 0
    var = x[:, 0] * 0

    for i in range(col):
        # data = x[0, :, i];
        data = x[:, i];
        mean = mean + data
        var += torch.square(data - mean / (i + 1))
    var = var / (col - 1)
    mean = mean / col

    mean.unsqueeze_(1)
    var.unsqueeze_(1)

    return self.alpha * (x - mean) / torch.sqrt(var + self.eps) + self.beta
```

# LayerNorm Optimization (Accuracy Analysis)

- Approach 1:
  - Fine-tune a pretrained BERT-base model on CoLA classification dataset
  - Replace all LayerNorm with LayerNorm\_v2 (keep the parameters)

```
# embedding layer
param = {}
# store the parameters
for name, data in new_model._modules['bert']._modules['embeddings']._modules['LayerNorm'].named_parameters():
    param[name] = data
# replace LayerNorm with LayerNorm_v2 (same parameters)
new_model._modules['bert']._modules['embeddings']._modules['LayerNorm'] = LayerNorm_v2(param['weight'], param['bias'])
```

- Evaluate with MCC metric (Matthews's correlation coefficient)
- The new model's predictions are stochastic (MCC = 0.1, baseline = 0.57)

# LayerNorm Optimization (Accuracy Analysis)

- Approach 2:
  - Replace LayerNorm in the pretrained model
    - With trainable parameters (weight, bias)
  - Fine-tune the new model
    - No loss drop on validation set
  - Still very low MCC score

```
Running Validation...  
Accuracy: 0.71  
Validation Loss: 0.61  
Validation took: 0:00:25
```

```
Running Validation...  
Accuracy: 0.71  
Validation Loss: 0.61  
Validation took: 0:00:25
```

```
Running Validation...  
Accuracy: 0.71  
Validation Loss: 0.60  
Validation took: 0:00:24
```

```
Running Validation...  
Accuracy: 0.71  
Validation Loss: 0.60  
Validation took: 0:00:25
```

# LayerNorm Optimization (Accuracy Analysis)

- Approach 3 (transformer from scratch):
  - Build the model layer by layer, LayerNorm vs LayerNorm\_v2
  - Extremely slow when training LayerNorm\_v2

Epoch Step:	1	Accumulation Step:	1	Loss:	2.95	Tokens / Sec:	132.1
Epoch Step:	41	Accumulation Step:	5	Loss:	2.73	Tokens / Sec:	134.8
Epoch Step:	81	Accumulation Step:	9	Loss:	2.78	Tokens / Sec:	134.5
Epoch Step:	121	Accumulation Step:	13	Loss:	2.83	Tokens / Sec:	134.3
Epoch Step:	161	Accumulation Step:	17	Loss:	2.74	Tokens / Sec:	132.7
Epoch Step:	201	Accumulation Step:	21	Loss:	2.68	Tokens / Sec:	133.9
Epoch Step:	241	Accumulation Step:	25	Loss:	2.64	Tokens / Sec:	134.9
Epoch Step:	281	Accumulation Step:	29	Loss:	2.86	Tokens / Sec:	135.0
Epoch Step:	321	Accumulation Step:	33	Loss:	2.40	Tokens / Sec:	134.2
Epoch Step:	361	Accumulation Step:	37	Loss:	2.64	Tokens / Sec:	135.3
Epoch Step:	401	Accumulation Step:	41	Loss:	2.50	Tokens / Sec:	133.9
Epoch Step:	441	Accumulation Step:	45	Loss:	2.62	Tokens / Sec:	133.7
Epoch Step:	481	Accumulation Step:	49	Loss:	2.66	Tokens / Sec:	134.6
Epoch Step:	521	Accumulation Step:	53	Loss:	2.52	Tokens / Sec:	134.8
Epoch Step:	561	Accumulation Step:	57	Loss:	2.61	Tokens / Sec:	134.9
Epoch Step:	601	Accumulation Step:	61	Loss:	2.58	Tokens / Sec:	133.4
Epoch Step:	641	Accumulation Step:	65	Loss:	2.38	Tokens / Sec:	135.1
Epoch Step:	681	Accumulation Step:	69	Loss:	2.49	Tokens / Sec:	135.5
Epoch Step:	721	Accumulation Step:	73	Loss:	2.39	Tokens / Sec:	132.5
Epoch Step:	761	Accumulation Step:	77	Loss:	2.64	Tokens / Sec:	133.6
Epoch Step:	801	Accumulation Step:	81	Loss:	2.20	Tokens / Sec:	134.5
Epoch Step:	841	Accumulation Step:	85	Loss:	2.48	Tokens / Sec:	133.0
Epoch Step:	881	Accumulation Step:	89	Loss:	2.54	Tokens / Sec:	136.0

Epoch Step:	1	Accumulation Step:	1	Loss:	2.05	Tokens / Sec:	2080.8
Epoch Step:	41	Accumulation Step:	5	Loss:	2.18	Tokens / Sec:	1715.0
Epoch Step:	81	Accumulation Step:	9	Loss:	2.30	Tokens / Sec:	1721.8
Epoch Step:	121	Accumulation Step:	13	Loss:	2.06	Tokens / Sec:	1725.1
Epoch Step:	161	Accumulation Step:	17	Loss:	2.00	Tokens / Sec:	1745.6
Epoch Step:	201	Accumulation Step:	21	Loss:	1.98	Tokens / Sec:	1717.5
Epoch Step:	241	Accumulation Step:	25	Loss:	1.60	Tokens / Sec:	1738.4
Epoch Step:	281	Accumulation Step:	29	Loss:	1.79	Tokens / Sec:	1733.4
Epoch Step:	321	Accumulation Step:	33	Loss:	1.96	Tokens / Sec:	1719.1
Epoch Step:	361	Accumulation Step:	37	Loss:	1.85	Tokens / Sec:	1703.4
Epoch Step:	401	Accumulation Step:	41	Loss:	2.07	Tokens / Sec:	1720.6
Epoch Step:	441	Accumulation Step:	45	Loss:	2.03	Tokens / Sec:	1731.8
Epoch Step:	481	Accumulation Step:	49	Loss:	1.88	Tokens / Sec:	1733.2
Epoch Step:	521	Accumulation Step:	53	Loss:	2.20	Tokens / Sec:	1726.6
Epoch Step:	561	Accumulation Step:	57	Loss:	1.61	Tokens / Sec:	1694.6
Epoch Step:	601	Accumulation Step:	61	Loss:	1.92	Tokens / Sec:	1735.5
Epoch Step:	641	Accumulation Step:	65	Loss:	2.01	Tokens / Sec:	1711.6
Epoch Step:	681	Accumulation Step:	69	Loss:	1.65	Tokens / Sec:	1714.4
Epoch Step:	721	Accumulation Step:	73	Loss:	1.69	Tokens / Sec:	1719.7
Epoch Step:	761	Accumulation Step:	77	Loss:	1.81	Tokens / Sec:	1732.5
Epoch Step:	801	Accumulation Step:	81	Loss:	1.49	Tokens / Sec:	1730.9
Epoch Step:	841	Accumulation Step:	85	Loss:	1.82	Tokens / Sec:	1717.7
Epoch Step:	881	Accumulation Step:	89	Loss:	1.63	Tokens / Sec:	1726.5

- Significantly higher loss (cannot finish the training due to colab GPU limitation)

# Bias & Concatenation

- Bias weight size: (feature\_size, 1)
  - One bias value for a column
- Embedding concatenation in MatMul:
  - Pass the head id to each head
  - Store the result vector to the position after concatenation based on head id

B1	B2	B3	B4

```
#ifndef SELF_ATTN_TEST
    vse32_v_f32m1(&o[i * dk + j], partial_sum, vl);
#else
    vse32_v_f32m1(&o[i * d_model + k * dk + j], partial_sum, vl);
#endif /* SELF_ATTN_TEST */
```