

Week 13

# New Instructions


- `vsetbl rd rs1`
- `vle<eew>bc rs1`
- `vfbmacc.vf vd, rs1, vs2`
- `vfbmacc.vv vd, vs2`

LLVM & Spike



# vsetbl


- set the broadcast length CSR
- CSR\_BL = rs1, rd = rs1
- 31-26: 1010000 (vsetvl: 1000000)



contain the length of  
broadcast vector b

## vle<eew>bc

- load vector to the broadcast buffer
- support all element widths
- use bl instead vl
- lumop: 11000 (unit-stride load: 00000)



# vfbmacc

- vfbmacc.vv vd, (vs1), vs2
  - vs1 unused
  - $vd[i, :] = bc\_vec[:] * vs2[i] + vd[i, :]$
- vfbmacc.vf vd, rs1, vs2
  - $vd[i, :] = bc\_vec[:] * vs2[i] + rs1$
  - to avoid initialization (slow)
  - $rs1 = 0$
- func6: 111001 (vfmacc: 101100)

# New MatMul Program

```
void fmatmul(float *c, const float *a, const float *b, const unsigned long int M,
             const unsigned long int N, const unsigned long int P) {
    const int REUSE_SIZE = 1;
    const int stride_a = 4 * N;
    const int stride_c = 4 * P;
    // We work on 64 elements of the matrix B at once
    // TODO: use 64 with LMUL=m2
    const unsigned long int block_size_p = 32;
    // block_size_m <= M, REUSE_SIZE * length of b_vec * 32 < VRF capacity
    const unsigned long int block_size_m = NR_LANES * REUSE_SIZE;

    for (unsigned long int p = 0; p < P; p += block_size_p) {
        // Find pointers to the submatrices
        const float *b_ = b + p;
        float *c_ = c + p;

        // Set the broadcast length
        const unsigned long int p_ = MIN(P - p, block_size_p);
        int tmp1 = 0;
        int tmp2 = 0;
        asm volatile("vsetbl %0, %1, %2 ::r"(tmp2), "r"(p_), "r"(tmp1));

        // Iterate over the rows
        for (unsigned long int m = 0; m < M; m += block_size_m) {
            // Set the vector length
            const unsigned long int m_ = MIN(M - m, block_size_m);
            asm volatile("vsetvli zero, %0, e32, m1, ta, ma ::r"(m_));

            // Find pointer to the submatrices
            const float *a_ = a + m * N;
            float *c__ = c_ + m * P;

            // First op with scalar zero
            // load vec_a
            asm volatile("vlse32.v v0, (%0), %1 ::r"(a_), "r"(stride_a));
            // load vec_b
            asm volatile("vle32bc.v v31, (%0) ::r"(b_));
            float t0 = 0; // First Operation, accumulated result is 0
            asm volatile("vfbmacc.vf v8, %0, v0 ::f"(t0));

            for (unsigned long int n = 1; n < N; n++) {
                // load vec_a
                asm volatile("vlse32.v v0, (%0), %1 ::r"(a_ + n), "r"(stride_a));
                // load vec_b
                asm volatile("vle32bc.v v31, (%0) ::r"(b_ + n * P));
                asm volatile("vfbmacc.vv v8, v31, v0");
            }

            asm volatile("vsetvli zero, %0, e32, m1, ta, ma ::r"(block_size_p * NR_LANES));
            asm volatile("vsse32.v v8, (%0), %1 ::r"(c__), "r"(stride_c));
        }
    }
}
```

# New MatMul Program

- broadcast length = 32
  - vlen = 4096 (length of one vreg)
  - vlen / #lanes / sizeof(float) = 32
  - bl = 64 (LMUL = 2)

```
void fmatmul(float *c, const float *a, const float *b, const unsigned long int M,
            const unsigned long int N, const unsigned long int P) {
    const int REUSE_SIZE = 1;
    const int stride_a = 4 * N;
    const int stride_c = 4 * P;
    // We work on 64 elements of the matrix B at once
    // TODO: use 64 with LMUL=4
    const unsigned long int block_size_p = 32;
    const unsigned long int m = M;
    // block_size_m = m; REUSE_SIZE * length of b_vec * 32 < VRF capacity
    const unsigned long int block_size_m = NR_LANES * REUSE_SIZE;

    for (unsigned long int p = 0; p < P; p += block_size_p) {
        // Find pointers to the submatrices
        const float *b_ = b + p;
        float *c_ = c + p;

        // Set the broadcast length
        const unsigned long int p_ = MIN(P - p, block_size_p);
        int tmp1 = 0;
        int tmp2 = 0;
        asm volatile("vsetbl %0, %1, %2" :: "r"(tmp2), "r"(p_), "r"(tmp1));

        // Iterate over the rows
        for (unsigned long int m = 0; m < M; m += block_size_m) {
            // Set the vector length
            const unsigned long int m_ = MIN(M - m, block_size_m);
            asm volatile("vsetvli zero, %0, e32, m1, ta, ma" :: "r"(m_));

            // Find pointer to the submatrices
            const float *a_ = a + m * N;
            float *c__ = c_ + m * P;

            // First op with scalar zero
            // load vec_a
            asm volatile("vlse32.v v0, (%0), %1" :: "r"(a_), "r"(stride_a));
            // load vec_b
            asm volatile("vle32bc.v v31, (%0)" :: "r"(b_));
            float t0 = 0; // First Operation, accumulated result is 0
            asm volatile("vfbmacc.vf v8, %0, v0" :: "f"(t0));

            for (unsigned long int n = 1; n < N; n++) {
                // load vec_a
                asm volatile("vlse32.v v0, (%0), %1" :: "r"(a_ + n), "r"(stride_a));
                // load vec_b
                asm volatile("vle32bc.v v31, (%0)" :: "r"(b_ + n * P));
                asm volatile("vfbmacc.vv v8, v31, v0");
            }

            asm volatile("vsetvli zero, %0, e32, m1, ta, ma" :: "r"(block_size_p * NR_LANES));
            asm volatile("vsse32.v v8, (%0), %1" :: "r"(c__), "r"(stride_c));
        }
    }
}
```

# New MatMul Program

- broadcast length = 32
  - vlen = 4096 (length of one vreg)
  - vlen / #lanes / sizeof(float) = 32
  - bl = 64 (LMUL = 2)
- reuse\_size
  - = number of accumulated registers
  - reuse\_size = 2, vd = v8 & v9
  - decrease memory access

```
void fmatmul(float *c, const float *a, const float *b, const unsigned long int M,
             const unsigned long int N, const unsigned long int P) {
    const int REUSE_SIZE = 1;
    const int stride_a = 4 * N;
    const int stride_c = 4 * P;
    // We work on 64 elements of the matrix B at once
    // TODO: use 64 with LMUL=m2
    const unsigned long int block_size_p = 32;
    // block_size_m <= M, REUSE_SIZE * length of b vec * 32 < VRF capacity
    const unsigned long int block_size_m = NR_LANES * REUSE_SIZE;

    for (unsigned long int p = 0; p < P; p += block_size_p) {
        // Find pointers to the submatrices
        const float *b_ = b + p;
        float *c_ = c + p;

        // Set the broadcast length
        const unsigned long int p_ = MIN(P - p, block_size_p);
        int tmp1 = 0;
        int tmp2 = 0;
        asm volatile("vsetbl %0, %1, %2" :: "r"(tmp2), "r"(p_), "r"(tmp1));

        // Iterate over the rows
        for (unsigned long int m = 0; m < M; m += block_size_m) {
            // Set the vector length
            const unsigned long int m_ = MIN(M - m, block_size_m);
            asm volatile("vsetvli zero, %0, e32, m1, ta, ma" :: "r"(m_));

            // Find pointer to the submatrices
            const float *a_ = a + m * N;
            float *c__ = c_ + m * P;

            // First op with scalar zero
            // load vec_a
            asm volatile("vlse32.v v0, (%0), %1" :: "r"(a_), "r"(stride_a));
            // load vec_b
            asm volatile("vle32bc.v v31, (%0)" :: "r"(b_));
            float t0 = 0; // First Operation, accumulated result is 0
            asm volatile("vfbmacc.vf v8, %0, v0" :: "f"(t0));

            for (unsigned long int n = 1; n < N; n++) {
                // load vec_a
                asm volatile("vlse32.v v0, (%0), %1" :: "r"(a_ + n), "r"(stride_a));
                // load vec_b
                asm volatile("vle32bc.v v31, (%0)" :: "r"(b_ + n * P));
                asm volatile("vfbmacc.vv v8, v31, v0");
            }

            asm volatile("vsetvli zero, %0, e32, m1, ta, ma" :: "r"(block_size_p * NR_LANES));
            asm volatile("vsse32.v v8, (%0), %1" :: "r"(c__), "r"(stride_c));
        }
    }
}
```



# New MatMul Program

- broadcast length = 32
  - vlen = 4096 (length of one vreg)
  - vlen / #lanes / sizeof(float) = 32
  - bl = 64 (LMUL = 2)
- reuse\_size
  - = number of accumulated registers
  - reuse\_size = 2, vd = v8 & v9
  - decrease memory access
- remove initialization

```
void fmatmul(float *c, const float *a, const float *b, const unsigned long int M,
             const unsigned long int N, const unsigned long int P) {
    const int REUSE_SIZE = 1;
    const int stride_a = 4 * N;
    const int stride_c = 4 * P;
    // We work on 64 elements of the matrix B at once
    // TODO: use 64 with LMUL=m2
    const unsigned long int block_size_p = 32;
    // block_size_m <= M, REUSE_SIZE * length of b vec * 32 < VRF capacity
    const unsigned long int block_size_m = NR_LANES * REUSE_SIZE;

    for (unsigned long int p = 0; p < P; p += block_size_p) {
        // Find pointers to the submatrices
        const float *b_ = b + p;
        float *c_ = c + p;

        // Set the broadcast length
        const unsigned long int p_ = MIN(P - p, block_size_p);
        int tmp1 = 0;
        int tmp2 = 0;
        asm volatile("vsetbl %0, %1, %2" :: "r"(tmp2), "r"(p_), "r"(tmp1));

        // Iterate over the rows
        for (unsigned long int m = 0; m < M; m += block_size_m) {
            // Set the vector length
            const unsigned long int m_ = MIN(M - m, block_size_m);
            asm volatile("vsetvli zero, %0, e32, m1, ta, ma" :: "r"(m_));

            // Find pointer to the submatrices
            const float *a_ = a + m * N;
            float *c__ = c_ + m * P;

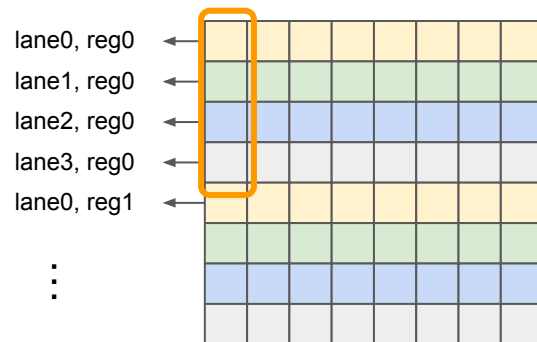
            // First op with scalar zero
            // load vec_a
            asm volatile("vlse32.v v0, (%0), %1" :: "r"(a_), "r"(stride_a));
            // load vec_b
            asm volatile("vle32bc.v v31, (%0)" :: "r"(b_));
            float t0 = 0; // First Operation, accumulated result is 0
            asm volatile("vfbmacc.vf v8, %0, v0" :: "f"(t0));

            for (unsigned long int n = 1; n < N; n++) {
                // load vec_a
                asm volatile("vlse32.v v0, (%0), %1" :: "r"(a_ + n), "r"(stride_a));
                // load vec_b
                asm volatile("vle32bc.v v31, (%0)" :: "r"(b_ + n * P));
                asm volatile("vfbmacc.vv v8, v31, v0");
            }

            asm volatile("vsetvli zero, %0, e32, m1, ta, ma" :: "r"(block_size_p * NR_LANES));
            asm volatile("vsse32.v v8, (%0), %1" :: "r"(c__), "r"(stride_c));
        }
    }
}
```




Result Matrix C



4 lanes,  
reuse\_size=2


# New store instruction

Result Matrix C  
stride store with baseAddr c



4 lanes,  
reuse\_size=2


# New store instruction



# New store instruction

Result Matrix C

stride store with baseAddr  
 $c + 2$



4 lanes,  
reuse\_size=2